

5118020-03 Operating Systems

Process

OSTEP Chapters 4 & 6

Shin Hong

Motivation

2



Process

--

ITP 30002
Operating System

2024-03-26

Process

3

- a running instance of a program
 - program vs. process
 - a kernel object that contains all information and resources given to the running instance
 - identified by a unique number (process ID)
- time-sharing of a CPU provides the illusion of many CPUs
 - concurrency vs. parallelism
 - policies: scheduling
 - mechanism: context switching

Constitution of Program Execution Context

4

- memory states
 - address space
- CPU states
 - registers: general-purpose and special-purpose
- I/O information

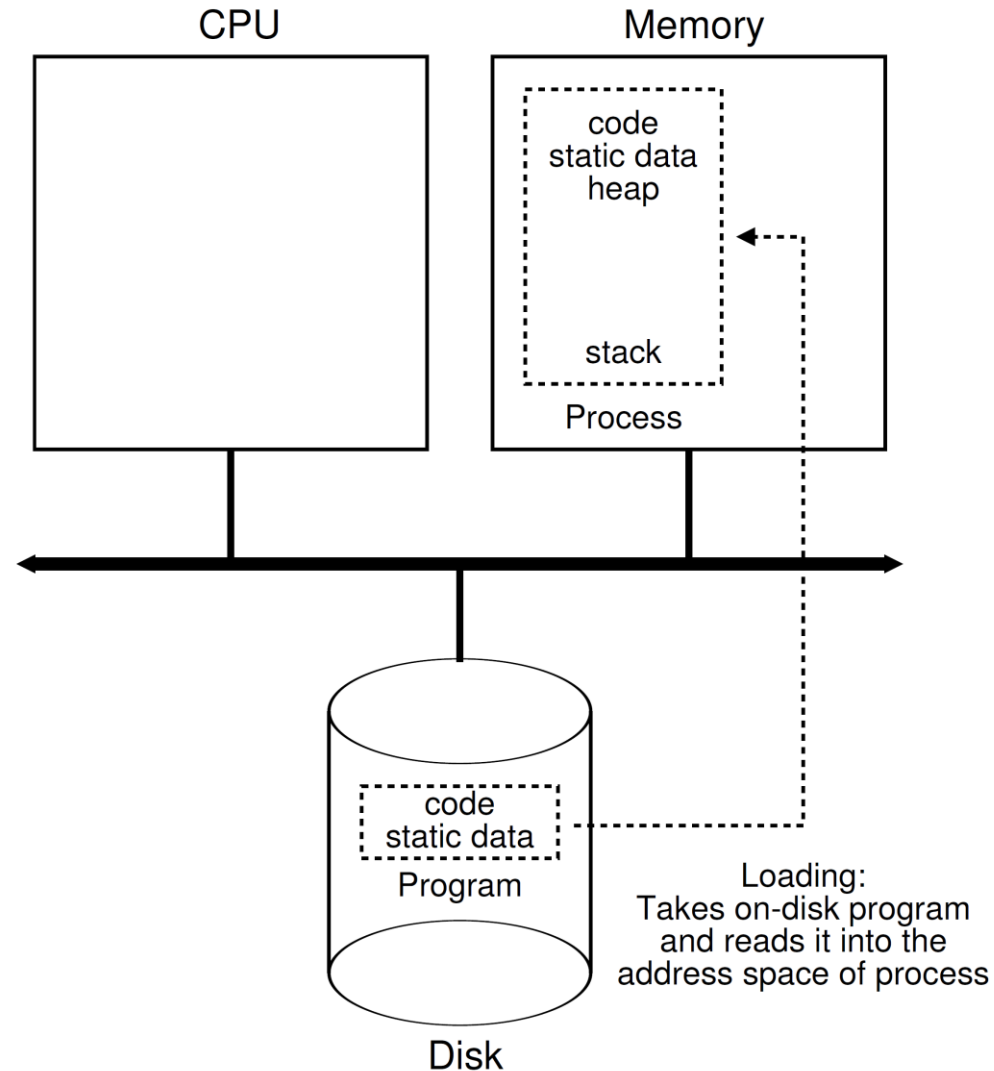
Process
--
ITP 30002
Operating System

2024-03-26

Life Cycle of a Process

5

- process creation
 - resource allocation
 - loading
 - eager manner
 - lazy manner



Process
--
ITP 30002
Operating System

2024-03-26

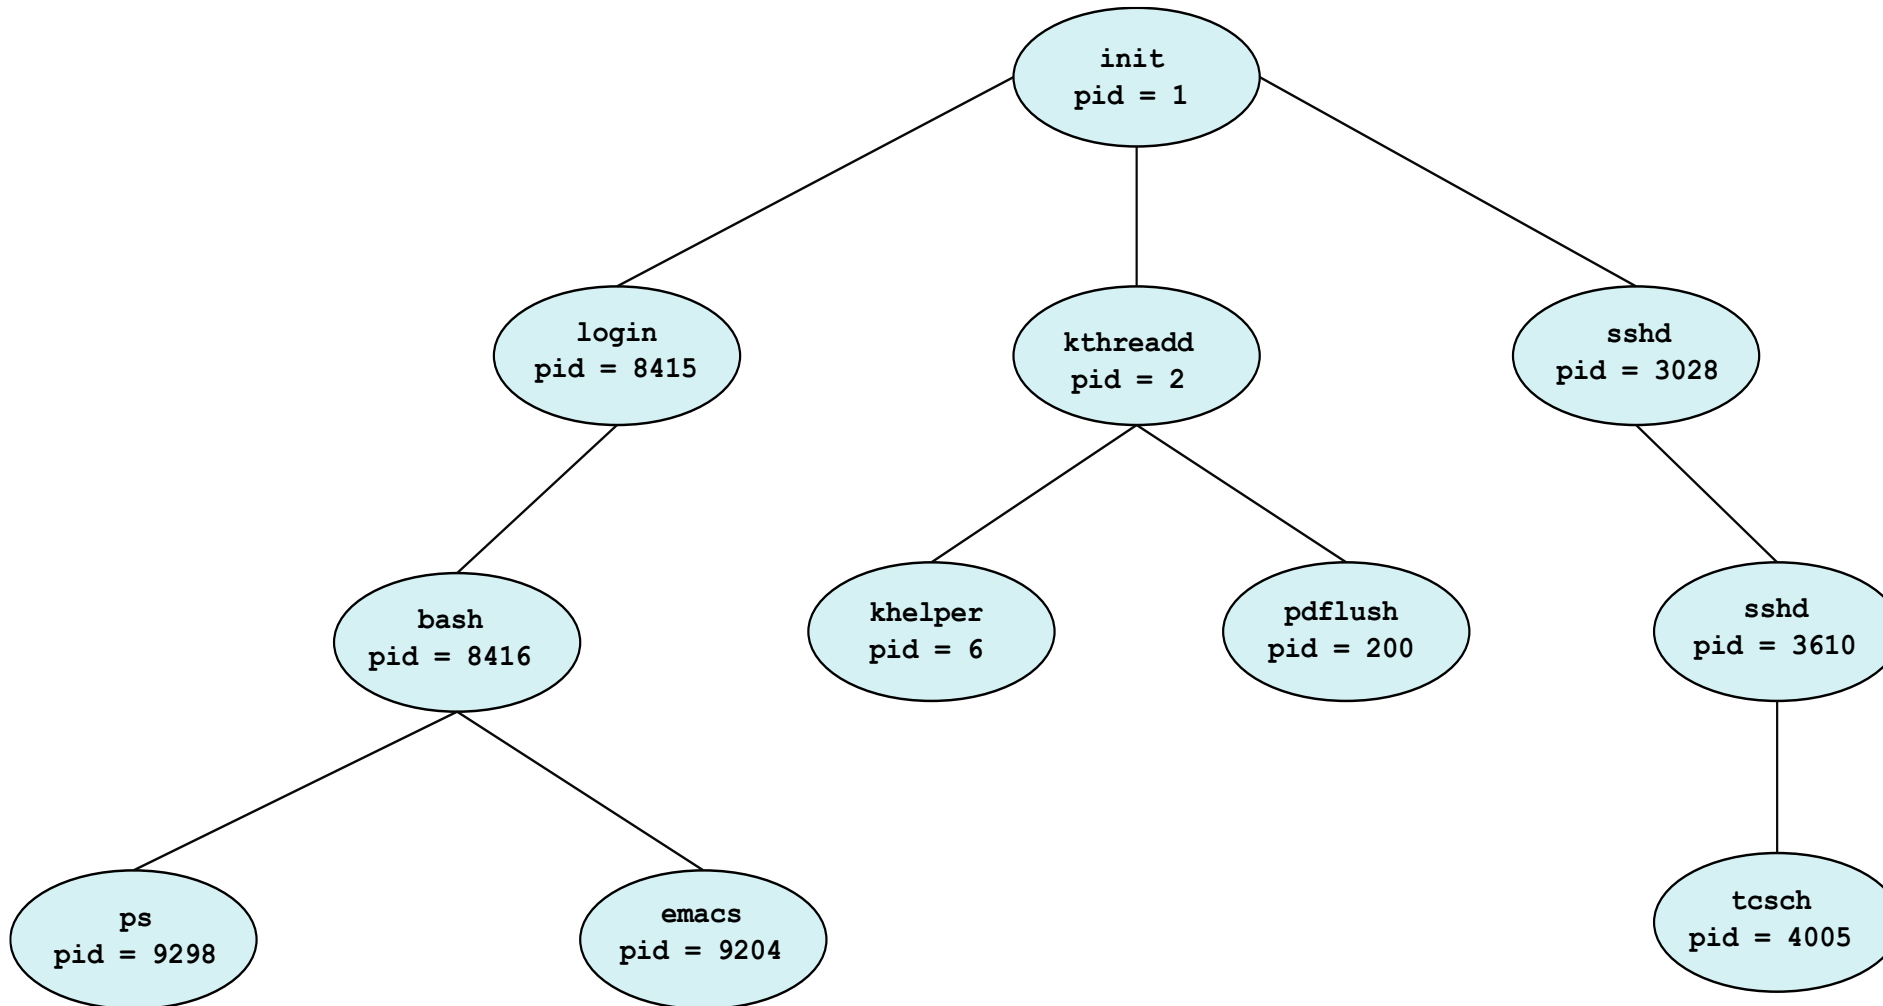
Process Creation

6

- A process is identified and managed via a process identifier (pid)
- A parent process can spawn a child process to delegate a subtask
 - A process can spawn multiple children processes
 - A parent process can run concurrently with its children processes
 - A child process, in turn create other processes, forming a tree of processes
 - A parent can wait until a child (or children) terminates
- A parent and its children can share resources
 - Children may share a subset of parent's resources
- Process in UNIX
 - a system call **fork()** system call creates a new process
 - a child process duplicates the memory of its parent

A Tree of Processes in Linux


7



Life Cycle of a Process

8

- Running state
 - hold a CPU and execute instructions
 - Ready state
 - can make a progress, but do not hold a CPU
 - Blocked state
 - cannot make a progress since it wait for a certain event/condition (i.e., I/O)
-
- CPU scheduler decides process state transitions



Example

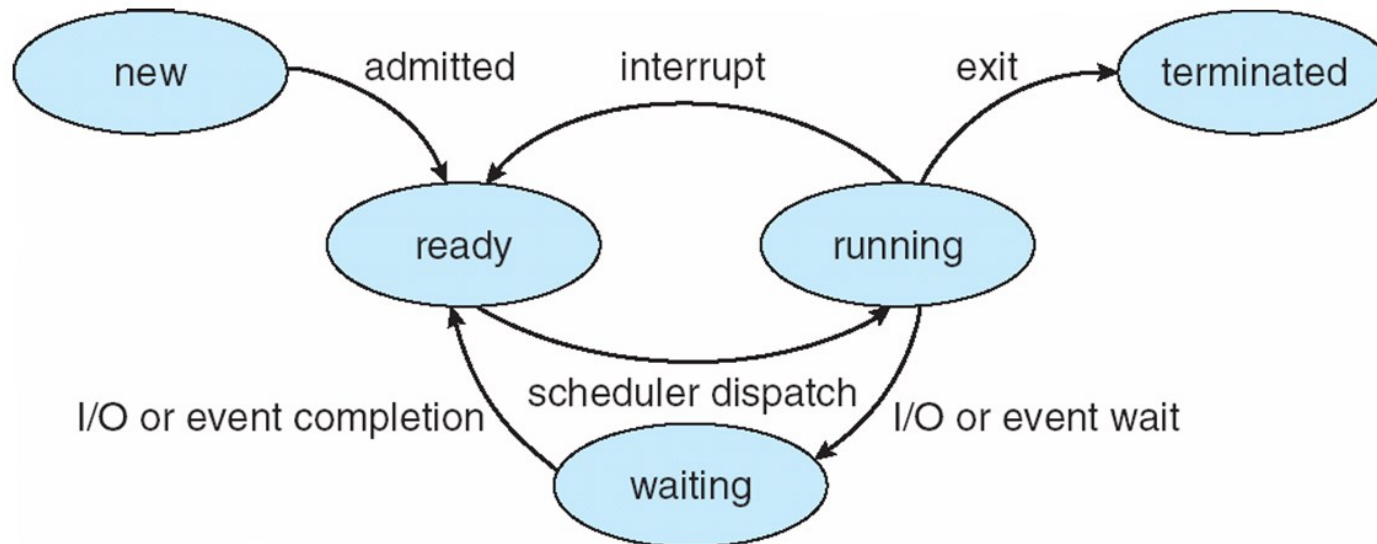
Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	
4	Running	Ready	Process ₀ now done
5	–	Running	
6	–	Running	
7	–	Running	
8	–	Running	Process ₁ now done

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	Process ₀ initiates I/O
4	Blocked	Running	Process ₀ is blocked,
5	Blocked	Running	so Process ₁ runs
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	Process ₁ now done
9	Running	–	
10	Running	–	Process ₀ now done

Another Representation of Process Life Cycle

10

- As a process executes, the process state changes
 - **new**: The process is being created
 - **ready**: The process is waiting to be assigned to a processor
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **terminated**: The process has finished execution



Process Termination

- When it finishes the given task, a process invokes the **exit** system call to delete itself
 - Given resources are reclaimed (deallocated)
 - Returns the exit status to the parent process (via **wait()**)
- A parent process can cut off a child process with **abort** system call
 - Examples
 - when a child exhausts all resources or time
 - when the task given to a child is not needed anymore
 - when the parent quits

Process Data Structure in Kernel: xv6 Example

12

- Called as Process Control Block
- Example of the xv6 kernel

```
// the registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    int eip;
    int esp;
    int ebx;
    int ecx;
    int edx;
    int esi;
    int edi;
    int ebp;
};

// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };
```

```
// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;                // Start of process memory
    uint sz;                  // Size of process memory
    char *kstack;             // Bottom of kernel stack
                                // for this process
    enum proc_state state;    // Process state
    int pid;                  // Process ID
    struct proc *parent;      // Parent process
    void *chan;               // If !zero, sleeping on chan
    int killed;               // If !zero, has been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;        // Current directory
    struct context context;    // Switch here to run process
    struct trapframe *tf;     // Trap frame for the
                                // current interrupt
};
```

Operating System

2024-03-26

Remainig Issues

13

- What's the overhead for virtualizing CPU?
- How to implement context switching?
- How to manage multiple processes?
- How a scheduler determines the next process to run?

Process
--
ITP 30002
Operating System

2024-03-26

Mechanisms for CPU Virtualization

14

- time-sharing
 - run one process for a while, and then switch to another one
 - issue
 - managing control
 - obtaining performance
- limited direct execution
 - run an application program directly on the CPU with some restriction
 - restriction
 - restricted memory accesses (H/W manipulation, resource allocation, etc.)
 - restricted instruction



Process
--
ITP 30002
Operating System

2024-03-26

Dual Mode Operation: User Mode & Kernel Mode

15

- Most contemporary processors provide dual mode operation
- User mode
 - for application program execution
 - restriction is enforced
 - a trap occurs when a process executes a restricted instruction under user mode
- Kernel model
 - for kernel execution
 - all privileged operations can be executed
- What if an application program needs to execute privileged operations?

Process
--
ITP 30002
Operating System

2024-03-26

System Call

16

- A mechanism for an application program to call the kernel to run to provide expected services
 - not possible to do this with the function-call mechanism
- To execute a system call, a program must execute a trap instruction
 - the operation of a trap instruction
 - change the mode into the kernel mode
 - store the current PC at a kernel stack
 - jump to a predefined program location for handling the trap
 - trap table
 - trap handler
 - each system call is identified by its unique number (i.e. system-call number)

Process

--

ITP 30002

Operating System

2024-03-26

Example of System Call Workflow

OS @ run (kernel mode)	Hardware	Program (user mode)
Create entry for process list Allocate memory for program Load program into memory Setup user stack with argv Fill kernel stack with reg/PC return-from-trap		
	restore regs (from kernel stack) move to user mode jump to main	Run main() ... Call system call trap into OS
	save regs (to kernel stack) move to kernel mode jump to trap handler	
Handle trap Do work of syscall return-from-trap		
	restore regs (from kernel stack) move to user mode jump to PC after trap	
		... return from main trap (via <code>exit()</code>)
Free memory of process Remove from process list		

Switching Between Processes

18

- How can OS regain control of the CPU when it is given to an application program?
- Natural chance: blocked operation
- Periodic scheduling
 - Cooperative approach
 - blocked operation
 - Preemptive scheduling approach
 - exploit a timer interrupt and its interrupt handler
 - requires HW support

Process
--
ITP 30002
Operating System

2024-03-26

Preemptive Scheduling

19

- scheduler: a kernel module that determines which process to dispatch at a chance (e.g., timer interrupt)
- context switch: a kernel module that is executed to switch processes running on a CPU
 - registered as a timer interrupt handler
 - steps
 - store the process status of a currently-running process to memory
 - CPU states
 - find the stored status of the next process from the memory
 - restore the stored status at the CPU
 - return the control back to the application program

Process
--
ITP 30002
Operating System

2024-03-26

Example

OS @ boot (kernel mode)	Hardware	
initialize trap table	remember addresses of... syscall handler timer handler	
start interrupt timer	start timer interrupt CPU in X ms	
OS @ run (kernel mode)	Hardware	Program (user mode)
		Process A
		...
	timer interrupt save regs(A) → k-stack(A) move to kernel mode jump to trap handler	
Handle the trap Call <code>switch()</code> routine save regs(A) → <code>proc_t(A)</code> restore regs(B) ← <code>proc_t(B)</code> switch to k-stack(B) return-from-trap (into B)		
	restore regs(B) ← k-stack(B) move to user mode jump to B's PC	
		Process B
		...

The xv6 Context Switch Code

21

```
1  # void swtch(struct context **old, struct context *new);
2  #
3  # Save current register context in old
4  # and then load register context from new.
5  .globl swtch
6  swtch:
7      # Save old registers
8      movl 4(%esp), %eax # put old ptr into eax
9      popl 0(%eax)      # save the old IP
10     movl %esp, 4(%eax) # and stack
11     movl %ebx, 8(%eax) # and other registers
12     movl %ecx, 12(%eax)
13     movl %edx, 16(%eax)
14     movl %esi, 20(%eax)
15     movl %edi, 24(%eax)
16     movl %ebp, 28(%eax)
17
18     # Load new registers
19     movl 4(%esp), %eax # put new ptr into eax
20     movl 28(%eax), %ebp # restore other registers
21     movl 24(%eax), %edi
22     movl 20(%eax), %esi
23     movl 16(%eax), %edx
24     movl 12(%eax), %ecx
25     movl 8(%eax), %ebx
26     movl 4(%eax), %esp # stack is switched here
27     pushl 0(%eax)      # return addr put in place
28     ret               # finally return into new ctxt
```

Process
--
ITP 30002
Operating System

2024-03-26