

5118020-03 Operating Systems

# Multi-level Feedback Queue

OSTEP Chapter 8

Shin Hong

# Multi-level Feedback Queue (MLFQ)

2

- classify processes into multiple levels wrt. their interactiveness
- run priority scheduling for processes across different levels, and run fair scheduling for processes of the same level
  - aims to achieve short response time of interactive processes, and also to achieve short turnaround time of computation-intensive processes
- predict next period of CPU-burst time of a process based on history




Multi-level  
Feedback Queue

5118020-03  
Operating Systems

2024-04-01

# MLFQ Mechanism

- Use multiple ready queues
  - each queue is assigned with a unique priority number
  - processes in the same queue have the same priority
- Scheduling algorithms: which process to dispatch next?
  - Rule 1. find the highest non-empty priority
  - Rule 2. select a process in a RR manner from the selected queue
- How to decide to the priority (queue) a process belongs?
  - assigned by the user
  - determined by the observed behaviors of runnable processes



Multi-level  
Feedback Queue

5118020-03  
Operating Systems

2024-04-01

# Change Priority Level of Process

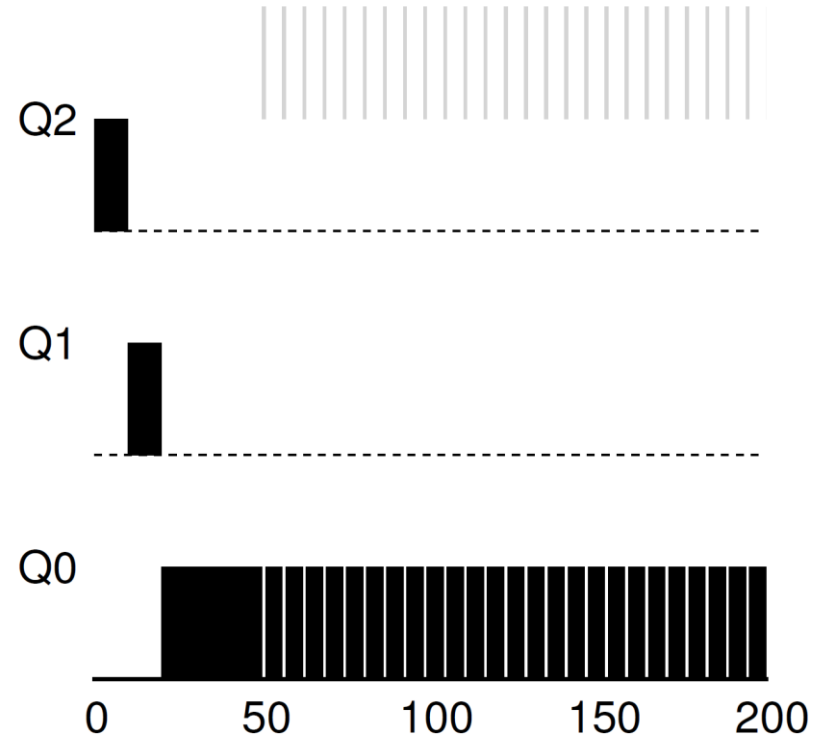

4

- Observations
  - an interactive process typically runs for a short period of time and quickly gets blocked (relinquishes the CPU)
  - a CPU-intensive process typically uses up a given CPU time and gets preempted
- Scheduling algorithm for controlling process priority
  - Rule 3. a new process is initially placed at the highest priority
  - Rule 4. a process is degraded to one level low priority if it uses up a time slice
  - Rule 5. a process stays at the same priority if it releases a CPU without preemption

Multi-level  
Feedback Queue

5118020-03  
Operating Systems  
2024-04-01

# Examples



5



Multi-level  
Feedback Queue

5118020-03  
Operating Systems

2024-04-01

# Problems of Priority Scheduling

6



- **Problems**

- non-interactive processes can be left out from scheduling if there are too many interactive ones; they suffer starvation
- a process has no chance to upgrade its priority even if its behavior changes



- **Solution:** periodically move the priorities of all processes to the top, every  $S$  time

Multi-level  
Feedback Queue

5118020-03  
Operating Systems  
2024-04-01

# Avoiding Gaming

7



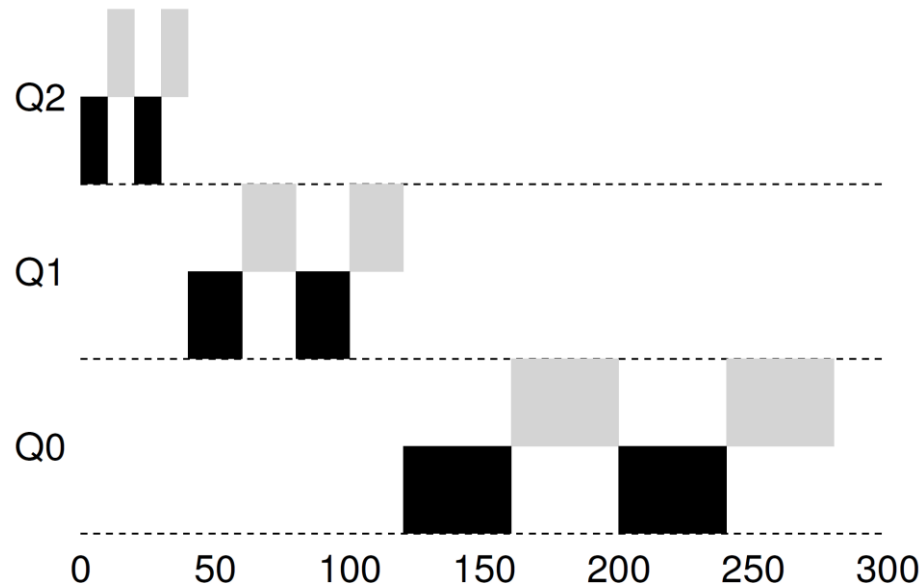
- **Problem:** A user can program to trick the scheduling by putting meaningless blocking operations to keep the process in a high priority queue
- **Solution:** if a process uses up a time allotment, move it one level down no matter how quickly it was to release the CPU (i.e., aging)

Multi-level  
Feedback Queue

5118020-03  
Operating Systems  
2024-04-01

# Parameter Tuning

8



- Performance will largely depend on parameters of scheduling policies
  - length of time slice, priority boosting frequency, etc.
  - Ex.: give a longer time slice to a process in a lower priority queue
  - Problem of *Voo-doo* constants
- Some systems use hints or commands from the user at scheduling

Multi-level  
Feedback Queue



5118020-03  
Operating Systems  
2024-04-01



# c.f. Danger of Priority Scheduling

9

- Priority scheduler strictly dispatches a ready process of the highest priority at a time
  - often used in real-time systems for strong completion time guarantee
- Under priority scheduling, multiple processes can be stuck (i.e., deadlock) if a process with a higher priority is waiting for a resource held by a process with a lower priority
  - E.g., *What really happened on Mars Rover PathFinder* by Mike Jones, Risks Digests, 1997



Multi-level  
Feedback Queue

5118020-03  
Operating Systems

2024-04-01